

✓ Lab Sheet 2

Implement and demonstrate Multiple Linear Regression for Brain Weights Prediction using sklearn Read the training data from a HeadBrain.CSV file.

<https://www.kaggle.com/datasets/jemishdonda/headbrain/data>

✓ Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

✓ Read Dataset

```
# Load the dataset
df = pd.read_csv('/kaggle/input/head-brain/headbrain.csv')
```

```
# Display the first few rows of the dataframe
print(df.head())
```

```
Gender  Age Range  Head Size(cm^3)  Brain Weight(grams)
0      1      1      4512      1530
1      1      1      3738      1297
2      1      1      4261      1335
3      1      1      3777      1282
4      1      1      4177      1590
```

```
df.shape
```

```
(237, 4)
```


```
df.tail()
```

```
Gender  Age Range  Head Size(cm^3)  Brain Weight(grams)
232     2      2      3214      1110
233     2      2      3394      1215
234     2      2      3233      1104
235     2      2      3352      1170
236     2      2      3391      1120
```

```
# Plot boxplot for each feature in the dataset
plt.figure(figsize=(14, 8))
sns.boxplot(data=df)
plt.title('Boxplot of Each Features in the HeadBrain Dataset', fontsize=16)
plt.xticks(rotation=45)
plt.show()
```







Boxplot of Each Features in the HeadBrain Dataset




```
# Plot histograms for each features
df.hist(figsize=(12, 8), bins=28)
plt.suptitle('Diatribution of Each Feature', fontsize=16)
plt.show()
```




Distribution of Each Feature




```
# Plot distribution plots with KDE for each feature
for column in df.columns:
    plt.figure(figsize=(8,4))
    sns.histplot(df[column], kde=True)
    plt.title(f'Distribution of {column}', fontsize=14)
    plt.xlabel(column)
    plt.ylabel('Frequency')
    plt.show()
```


 /opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed with pd.option_context('mode.use_inf_as_na', True):




/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed with pd.option_context('mode.use_inf_as_na', True):

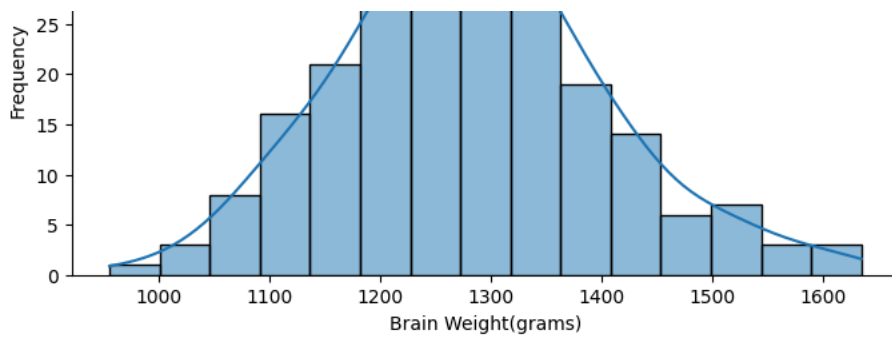


/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed with pd.option_context('mode.use_inf_as_na', True):



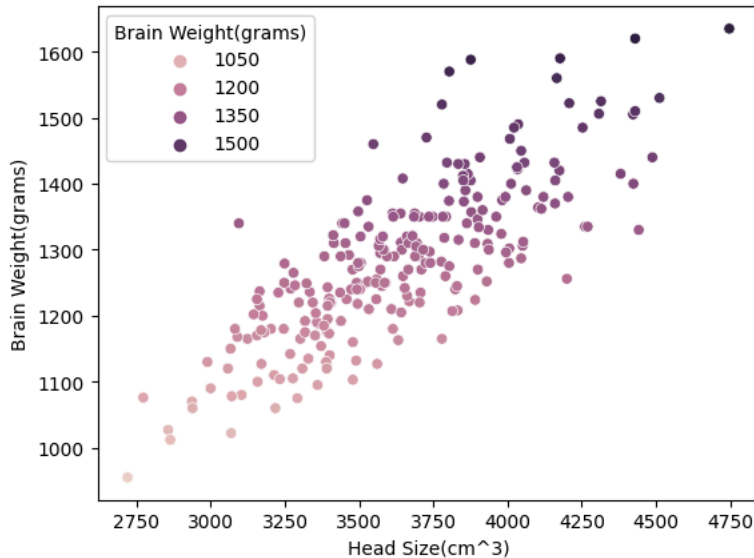
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed with pd.option_context('mode.use_inf_as_na', True):





```
sns.scatterplot(x='Head Size(cm^3)', y='Brain Weight(grams)', hue='Brain Weight(grams)', data=df)
```

```
<Axes: xlabel='Head Size(cm^3)', ylabel='Brain Weight(grams)'\>
```



```
#Calculate the correlation matrix
```

```
##df.corr(): Computes the correlation matrix for the DataFrame, showing the correlation coefficients between variables.
```

```
tc=df.corr()
```


```
tc
```




	Gender	Age Range	Head Size(cm^3)	Brain Weight(grams)
Gender	1.000000	-0.088652	-0.514050	-0.465266
Age Range	-0.088652	1.000000	-0.105428	-0.169438
Head Size(cm^3)	-0.514050	-0.105428	1.000000	0.799570
Brain Weight(grams)	-0.465266	-0.169438	0.799570	1.000000

```
plt.figure(figsize=(10,5))
```

```
sns.heatmap(tc)
```

 <Axes: >


```
#Visualize the correlation matrix using a heatmap
sns.heatmap(): Creates a heatmap to visualize the correlation matrix.
##annot=True: Displays the correlation coefficient values in the heatmap.
##cmap='coolwarm': Sets the color map for the heatmap.
plt.figure(figsize=(10,5))
tc=df.corr().round(2)
sns.heatmap(tc,annot=True,cmap='coolwarm')
```


 <Axes: >


```
# Identify and list the features with the highest positive and negative correlation with the target variable
target_correlation = tc['Brain Weight(grams)'].sort_values(ascending=False)
```

```
top_positive_correlations = target_correlation.head()
```

```
top_negative_correlations = target_correlation.tail()
```

```
print(" Positive Correlations:\n", top_positive_correlations)
print("\n Negative Correlations:\n", top_negative_correlations)
```

 Positive Correlations:
Brain Weight(grams) 1.00
Head Size(cm³) 0.80
Age Range -0.17
Gender -0.47
Name: Brain Weight(grams), dtype: float64

Negative Correlations:
Brain Weight(grams) 1.00

```

Head Size(cm^3)      0.80
Age Range            -0.17
Gender               -0.47
Name: Brain Weight(grams), dtype: float64

```

✓ Prepare the data for training the Linear Regression model

```

#Select the features and the target variable
#df.drop(): Removes the target variable 'Brain Weight' from the features DataFrame X.
#df['Brain Weight(grams)']: Selects the target variable y.

```

```

# Select the features and the target variable
X=df.drop('Brain Weight(grams)',axis=1)
y=df['Brain Weight(grams)']

```

df.drop(): Removes the target variable 'MEDV' from the features DataFrame X.

df['Brain Weight(grams)']: Selects the target variable y.

```

#Split the dataset into training and testing sets

```

```

##train_test_split(): Splits the data into training and testing sets.
##test_size=0.2: Allocates 20% of the data for testing and 80% for training.
##random_state=42: Ensures reproducibility of the split.
from sklearn.model_selection import train_test_split

```

```

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

train_test_split(): Splits the data into training and testing sets. **test_size=0.2:** Allocates 20% of the data for testing and 80% for training.

random_state=42: Ensures reproducibility of the split.

✓ Standardize the feature variables

```

from sklearn.preprocessing import StandardScaler

```

```

# Standardize the feature variables
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

```

StandardScaler(): Standardizes features by removing the mean and scaling to unit variance. **fit_transform():** Fits the scaler to the training data and transforms it. ***transform():** *Transforms the testing data using the already fitted scaler.

✓ Train a Linear Regression model using the training data

```

from sklearn.linear_model import LinearRegression
model=LinearRegression()
model.fit(X_train,y_train)

```

```

LinearRegression
LinearRegression()

```

In Training,62% Score

```

model.score(X_train,y_train)*100

```

```

62.64565800251868

```

```

# Display the model's coefficients and intercept
print("Coefficients:", model.coef_)

```